

# Dataset Provenance, Clinical Semantics, and Machine Learning Suitability

A Comprehensive Evaluation for the PneumoNet Multi-Stage Deep Learning Framework

Author: Medical AI Research Group | Project: PneumoNet Thesis Methodology | Domain: Computer-Aided Thoracic Radiography

## Executive Summary & System-Level Integration

This document presents a rigorous academic and technical analysis of three foundational chest radiograph datasets: **NIH ChestX-ray14**, **Curated Chest X-Ray Dataset for COVID-19**, and the **TBX11K Benchmark Dataset**. The evaluation scrutinizes dataset provenance, clinical semantics, label extraction mechanics, mathematical task formalizations, data leakage vectors, pre-processing trade-offs, and dataset-level limitations to establish a scientifically sound methodology for the **PneumoNet** deep learning pipeline.


## 1. Comparative Dataset Architectural Matrix

| Attribute / Level           | NIH ChestX-ray14                                        | Curated COVID-19 Dataset                                    | TBX11K Benchmark                                              |
|-----------------------------|---------------------------------------------------------|-------------------------------------------------------------|---------------------------------------------------------------|
| Primary Domain / Scope      | Broad thoracic pathologies (14 disease categories)      | Viral & bacterial pulmonary infections                      | Tuberculosis screening, classification & localization         |
| Provenance & Origin         | NIH Clinical Center (Single-institution hospital scale) | 15 distinct aggregated public image repositories            | Primary clinical collection across partner diagnostic centers |
| Scale (Images / Patients)   | 112,120 images / 30,805 unique patients                 | 9,208 images / Patient count undefined                      | 11,200 images / De-identified clinical cases                  |
| Annotation Level            | Weakly supervised image-level labels                    | Image-level classification labels                           | Bounding box spatial coordinates (COCO/Pascal VOC XML)        |
| Label Extraction Method     | Automatic NLP mining of unformatted radiology reports   | Consensus compilation from source datasets                  | Expert radiologist manual annotation                          |
| Labeling Noise / Error      | Modest error (~10–15% noise depending on class)         | Low label noise, high source heterogeneity noise            | Minimal label noise; expert ground truth                      |
| Target Task Formalization   | Multi-label binary classification (14 sigmoid outputs)  | Multi-class single-label classification (4 softmax outputs) | Multi-task: Multi-class classification + Object Detection     |
| Native Spatial Resolution   | 1024 × 1024 pixels                                      | Varied (224 × 224 up to 2000 × 2000+)                       | 512 × 512 pixels                                              |
| Primary Data Leakage Vector | Inter-view patient duplication across splits            | Image duplication across merged source repositories         | Over-represented healthy/sick-non-TB background classes       |
| Primary PneumoNet Role      | Primary Pre-training / Multi-label Backbone Model       | Zero-Shot / Domain Generalization Test Set                  | Ground-Truth Localization & Small-Lesion Benchmark            |

## 2. NIH ChestX-ray14: In-Depth Evaluation

### 2.1 Dataset Provenance & Collection History

Introduced by Wang et al. (CVPR 2017) as *ChestX-ray8* and expanded to *ChestX-ray14*, this dataset contains 112,120 frontal (PA/AP) chest X-rays from 30,805 unique patients admitted to the NIH Clinical Center between 1992 and 2015. It was created to provide the computer vision community with a massive, de-identified clinical radiography dataset to benchmark deep learning models.



### 2.2 Clinical Meaning & Disease Taxonomy

The dataset tracks 14 distinct thoracic pathologies:

- **Atelectasis:** Partial or complete collapse of a lung lobe or entire lung tissue.
- **Cardiomegaly:** Enlarged cardiac silhouette (cardiothoracic ratio > 0.50).
- **Effusion:** Pathological fluid accumulation in the pleural cavity.
- **Infiltration:** Dense substance (pus, blood, protein) trapped in lung parenchymal cells.
- **Mass:** Lesion greater than 30 mm in diameter, raising malignancy concerns.
- **Nodule:** Discrete, well-defined opacity up to 30 mm in diameter.
- **Pneumonia:** Acute inflammatory exudate filling alveoli due to infection.
- **Pneumothorax:** Trapped air within the pleural space causing lung displacement.
- **Consolidation:** Replacement of alveolar air with fluid, inflammatory exudate, or tissue.
- **Edema:** Fluid accumulation in tissue/air spaces, typically secondary to heart failure.
- **Emphysema:** Permanent enlargement and destruction of alveolar walls.
- **Fibrosis:** Chronic scarring and stiffening of pulmonary interstitium.
- **Pleural Thickening:** Fibrotic scarring or calcification of the pleural membrane.
- **Hernia:** Visceral displacement through anatomical defects (e.g., diaphragmatic).

### 2.3 Annotation Methodology & NLP Mining

Labels were generated automatically by running Natural Language Processing (NLP) tools—specifically **DNorm** and **MetaMap**—over raw unstructured radiology reports.

- **Labeling Strategy:** Extracted concepts were mapped to Medical Subject Headings (MeSH).
- **Ground-Truth Nature:** Labels are *weakly supervised* and noisy. The NLP extraction accuracy is estimated at ~90%. False positives stem from negated findings or historical comparisons in the report, while false negatives occur when radiologists omit implicit secondary findings.
- **Clinical Nuance ("No Finding"):** "No Finding" indicates the absence of the 14 targeted pathologies in the report. It does *not* guarantee a completely healthy patient, as unlisted pathologies (e.g., rib fractures, surgical clips, scoliosis) may still be present.

## 2.4 Structural Specification & Label Multiplicity

Mathematical Structure: Multi-label binary classification.

Label Vector: Each image  $x_i$  is mapped to a label vector  $y_i \in \{0, 1\}^{14}$ .

Co-occurrence: A single image can exhibit multiple positive flags simultaneously (e.g.,  $y_i = [Effusion=1, Cardiomegaly=1, Atelectasis=1, \dots]$ ).

## 2.5 Machine Learning Suitability

Ideal for training deep Convolutional Neural Networks (ResNet-50, DenseNet-121, EfficientNet-B4) using multi-label loss functions. Loss Function Formulation includes Binary Cross-Entropy (BCE) with class weighting, or Focal Loss to handle severe background sparsity:

$$\mathcal{L}_{BCE} = -\frac{1}{N} \sum_{i=1}^N \sum_{c=1}^{14} [w_c \cdot y_{i,c} \log(\hat{y}_{i,c}) + (1 - y_{i,c}) \log(1 - \hat{y}_{i,c})]$$


## 2.6 Pre-Processing & Spatial Trade-Offs

- **Native vs. Downsampled:** Native resolution is  $1024 \times 1024 \times 8$ -bit PNG.
- **Resizing Downside:** Standard resizing to  $224 \times 224$  pixels reduces computational overhead, but bilinear or bicubic downsampling degrades subtle high-frequency features (e.g., micro-nodules, fine fibrotic striations, subtle pneumothoraces).
- **Mitigation:** Pre-processing should use high-resolution training ( $512 \times 512$ ), Contrast Limited Adaptive Histogram Equalization (CLAHE), and image normalization tuned to ImageNet or dataset-specific statistics:

$$\mu = [0.485, 0.456, 0.406], \quad \sigma = [0.229, 0.224, 0.225]$$


## 2.7 Data Leakage & Validation Pitfalls

Because patients often undergo multiple follow-up X-rays over time, random image-level splitting causes severe data leakage. Splits **must** be grouped strictly by Patient ID. All scans from a single patient must belong exclusively to either the Training, Validation, or Testing set.



## 2.8 Class Imbalance & Metric Selection

Class frequencies are heavily skewed (e.g., Infiltration ~19,870 cases vs. Hernia ~227 cases).



### Metric Selection Rule

Standard accuracy is misleading because predicting all zeros yields high nominal accuracy. Required Evaluation Metrics include Per-class Area Under the Receiver Operating Characteristic Curve (AUROC), Area Under the Precision-Recall Curve (AUPRC), Macro-F1 score, and Sensitivity/Recall at fixed Specificity (95%).

## 2.9 Role in PneumoNet & Academic Justification


**PneumoNet Role:** Serves as the primary pre-training representation backbone.

**Academic Justification:** Pre-training on NIH ChestX-ray14 transfers domain-specific radiologic features far better than natural ImageNet features. It allows the model to learn universal chest representations before fine-tuning on specialized downstream tasks.

# 3. Curated COVID-19 Dataset: In-Depth Evaluation

## 3.1 Dataset Provenance & Curation Logic

Created by Unais Sait et al. during the COVID-19 pandemic, this dataset aggregates chest radiographs from 15 independent public repositories (including Cohen's COVID-19 image data collection, Italian Society of Medical and Interventional Radiology, Kaggle Chest X-Ray Images, etc.).



## 3.2 Clinical Meaning & Granular Breakdown

This dataset focuses on infectious lung diseases, categorizing images into four distinct classes:

|                     |             |       |
|---------------------|-------------|-------|
| Normal              | <div></div> | 3,270 |
| Bacterial Pneumonia | <div></div> | 3,001 |
| Viral Pneumonia     | <div></div> | 1,656 |
| COVID-19            | <div></div> | 1,281 |


- **Normal (n = 3,270):** Clear lung fields without focal opacities or consolidations.
- **COVID-19 (n = 1,281):** SARS-CoV-2 infection, typically showing bilateral peripheral ground-glass opacities (GGO) and lower-lobe consolidation.
- **Viral Pneumonia (Non-COVID) (n = 1,656):** Non-SARS-CoV-2 viral infections (e.g., Influenza, RSV), characterized by diffuse interstitial patterns and bronchial wall thickening.
- **Bacterial Pneumonia (n = 3,001):** Typically presents with localized, dense lobar consolidation.

## 3.3 Multi-Class Differential Diagnosis Significance

Distinguishing viral from bacterial pneumonia—and isolating COVID-19 from non-COVID viral etiologies—is a key diagnostic challenge. Including all four classes forces the network to learn genuine diagnostic markers rather than just binary "normal vs. abnormal" heuristics.

### 3.4 Multi-Source Aggregation & Domain Bias Risks

Because the images originate from 15 distinct sources, there is a risk of **Domain / Shortcut Bias**. If Source A supplies mostly COVID-19 cases and Source B supplies mostly Normal cases, a deep neural network may learn hardware artifacts (e.g., text overlays, background contrast, window margins, corner markers) rather than clinical pathology.



### 3.5 Deduplication Pipeline

The authors applied structural similarity metrics (SSIM) and perceptual hashing to remove cross-repository duplicate images, reducing the initial pool down to 9,208 high-quality images. This prevents identical images from appearing in both training and test splits.

### 3.6 Machine Learning Formalization & Loss Architecture

Task Type: 4-class single-label categorical classification mapped to one-hot encoded vector  $y_i \in \{0, 1\}^4$ .

Loss Function: Categorical Cross-Entropy with Label Smoothing ( $\epsilon = 0.1$ ) to account for potential cross-source annotation variations:

$$\mathcal{L}_{CE} = - \sum_{c=1}^4 y'_{i,c} \log(\hat{y}_{i,c}), \quad \text{where } y'_{i,c} = (1 - \epsilon) y_{i,c} + \frac{\epsilon}{4}$$

### 3.7 Pre-Processing Protocols

- **Normalization:** Heavy spatial cropping to remove extra-thoracic borders, text labels, and scanner markers.
- **Augmentation:** Random affine transformations (rotations  $\pm 10^\circ$ , translation  $\pm 5\%$ , zoom  $\pm 5\%$ ) and intensity adjustments. Cutout/Random Erasing should be used cautiously to avoid destroying localized ground-glass opacities.

### 3.8 Role in PneumoNet & Academic Justification


**PneumoNet Role:** Serves as an Out-of-Distribution (OOD) / External Generalization Test Set.

**Academic Justification:** Evaluating PneumoNet on an aggregated, multi-source dataset demonstrates model robustness across diverse imaging systems and patient populations, proving the system works beyond its primary training set.

## 4. TBX11K Benchmark: In-Depth Evaluation


### 4.1 Dataset Provenance & Structural Paradigm Shift

Introduced by Liu et al. (CVPR 2020) to address the lack of large-scale, annotated Tuberculosis (TB) benchmarks. While legacy TB datasets contained only a few hundred images (e.g., Montgomery: 138, Shenzhen: 662), TBX11K provides **11,200 high-resolution images** (512 × 512) along with spatial bounding-box annotations.



## 4.2 Clinical Category Semantics & Release Taxonomy

The released dataset distinguishes specific TB manifestations alongside non-TB categories:



- **Healthy (n = 5,000):** Normal chest radiographs.
- **Sick but Non-TB (n = 5,000):** Radiographs showing non-TB pulmonary conditions (e.g., asthma, pneumonia, cancer). Including this class prevents the network from oversimplifying the task into "TB vs. non-TB."
- **Active TB (n = 924):** Manifests as active lesions, cavitary opacities, apical infiltrates, and fibronodular shadows.
- **Latent TB (n = 212):** Inactive or controlled infection, typically presenting as calcified Ghon foci or stable fibrotic scars.
- **Active & Latent TB Co-existence (n = 54):** Displays both active lesions and old calcified scars.
- **Uncertain TB (n = 10):** Radiographic findings suggesting TB where clinical history cannot definitively confirm active versus latent status.

## 4.3 Bounding Box Annotations & Detection Paradigms

TBX11K provides bounding box annotations for affected lung regions in Pascal VOC XML and MS COCO JSON formats:

$$\mathbf{b}_i = [x_{\min}, y_{\min}, x_{\max}, y_{\max}, \text{class\_id}]$$

This enables the training of object detection and instance segmentation models (e.g., Faster R-CNN, RetinaNet, YOLOv8/v10, FCOS) alongside standard classification architectures.



## 4.4 Spatial Resolution & Small-Lesion Integrity

Images are provided natively at 512 × 512 resolution. Maintaining this spatial resolution is critical because tuberculosis features—such as subtle apical cavitary lesions, miliary nodules, and micronodular infiltrates—often disappear when downsampled to 224 × 224.

## 4.5 Standardized Splits & Challenge Protocols

The dataset includes a standardized benchmark split: **Training Set:** 6,600 images | **Validation Set:** 1,800 images | **Testing Set:** 2,800 images (Ground-truth labels preserved for server evaluation). To avoid validation leakage, researchers should evaluate on the official validation split or follow standard split protocols when testing locally.

## 4.6 Machine Learning Suitability & Multi-Task Losses


TBX11K supports Multi-Task Learning (MTL), combining image classification and region detection losses:

$$\mathcal{L}_{\text{Total}} = \mathcal{L}_{\text{cls}}(\hat{y}, y) + \lambda \sum_i \mathcal{L}_{\text{box}}(\hat{b}_i, b_i)$$

Where  $\mathcal{L}_{\text{cls}}$  is Cross-Entropy/Focal Loss,  $\mathcal{L}_{\text{box}}$  is Complete IoU (CloU) Loss, and  $\lambda$  is a balancing hyperparameter.

#### 4.7 Validating Grad-CAM Heatmaps

A major advantage of TBX11K is that its bounding boxes allow for objective validation of explainable AI (XAI) outputs like Grad-CAM. Instead of qualitatively assessing heatmaps, model attention can be quantified by comparing heatmaps against expert bounding boxes using metrics like Energy-based Pointing Game Accuracy or Intersection over Union (IoU).




#### 4.8 Role in PneumoNet & Academic Justification

**PneumoNet Role:** Serves as the dedicated TB Detection, Localization, and XAI Validation Subsystem.

**Academic Justification:** TBX11K allows PneumoNet to move beyond simple image-level predictions by pin-pointing specific lesions and rigorously verifying that its visual explanations align with expert annotations.

## 5. Comprehensive Cross-Dataset Analysis for Thesis

### 5.1 Methodological Comparison



### 5.2 Methodological Rigor & Data Leakage Prevention

To ensure valid performance estimates, the PneumoNet pipeline enforces strict data splitting rules across all three datasets:

- NIH ChestX-ray14:** Grouped strictly by Patient ID to prevent scans from the same patient from appearing in both training and test sets.
- Curated COVID-19 Dataset:** Screened using perceptual hashing to prevent cross-source duplicates from inflating performance during external validation.
- TBX11K Dataset:** Split using official benchmark assignments to ensure fair comparison with published literature.

### 5.3 PneumoNet Strategic Pipeline Integration

By combining these three datasets, PneumoNet addresses key challenges in medical AI:

- **NIH ChestX-ray14** provides the broad feature representation needed to analyze general chest X-rays.
- **Curated COVID-19 Dataset** validates that the model generalizes across different hospital systems, imaging equipment, and viral etiologies.
- **TBX11K Benchmark** adds precise spatial localization and allows visual explanations (Grad-CAM) to be quantitatively validated against expert bounding boxes.

This three-tiered approach provides a robust, clinically relevant framework for automated chest radiograph analysis.